

# Lecture 3: Introduction to Convolutional Neural Networks

Rafael Martínez-Galarza, Pavlos Protopapas  
Harvard-Smithsonian Center for Astrophysics

CENTER FOR

**ASTROPHYSICS**


HARVARD & SMITHSONIAN

# This lecture:

- The need for computer vision
- The architecture of a CNN
- Key features of CNNs
- ReLU activation function

# Convolutional Neural Networks

# A fully connected neural net



Each neuron in hidden layer 1 is connected to all neurons in the previous (input) layer and to all neurons in the following (hidden2) layer.

Very good for regression, for learning complex functions, but...

# Finding structure in images



Nearby pixels are more strongly related than distant ones.

Objects are built up out of smaller parts.

A fully connected network would be very inefficient at learning patterns here.


Instead, we can learn from nature.

# Orientation invariance






# Vision in nature



Cells in the visual cortex are not sensitive to the entire visual field  
First neurons connected to the retina specialize in detecting edges  
Neurons in further layers identify other types of patterns

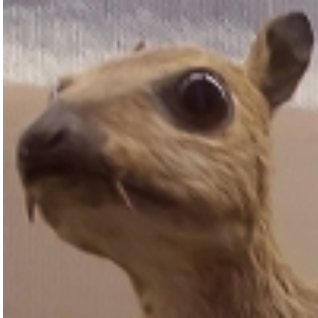

# “Convolution” operation





# Feature extraction

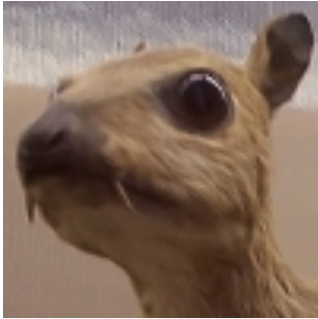
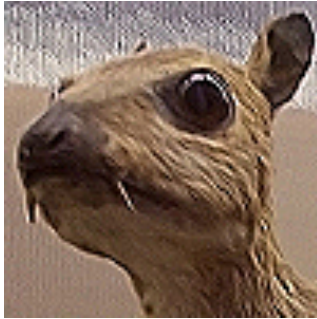
*Edge detection*


$$* \begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix} =$$


Kernel

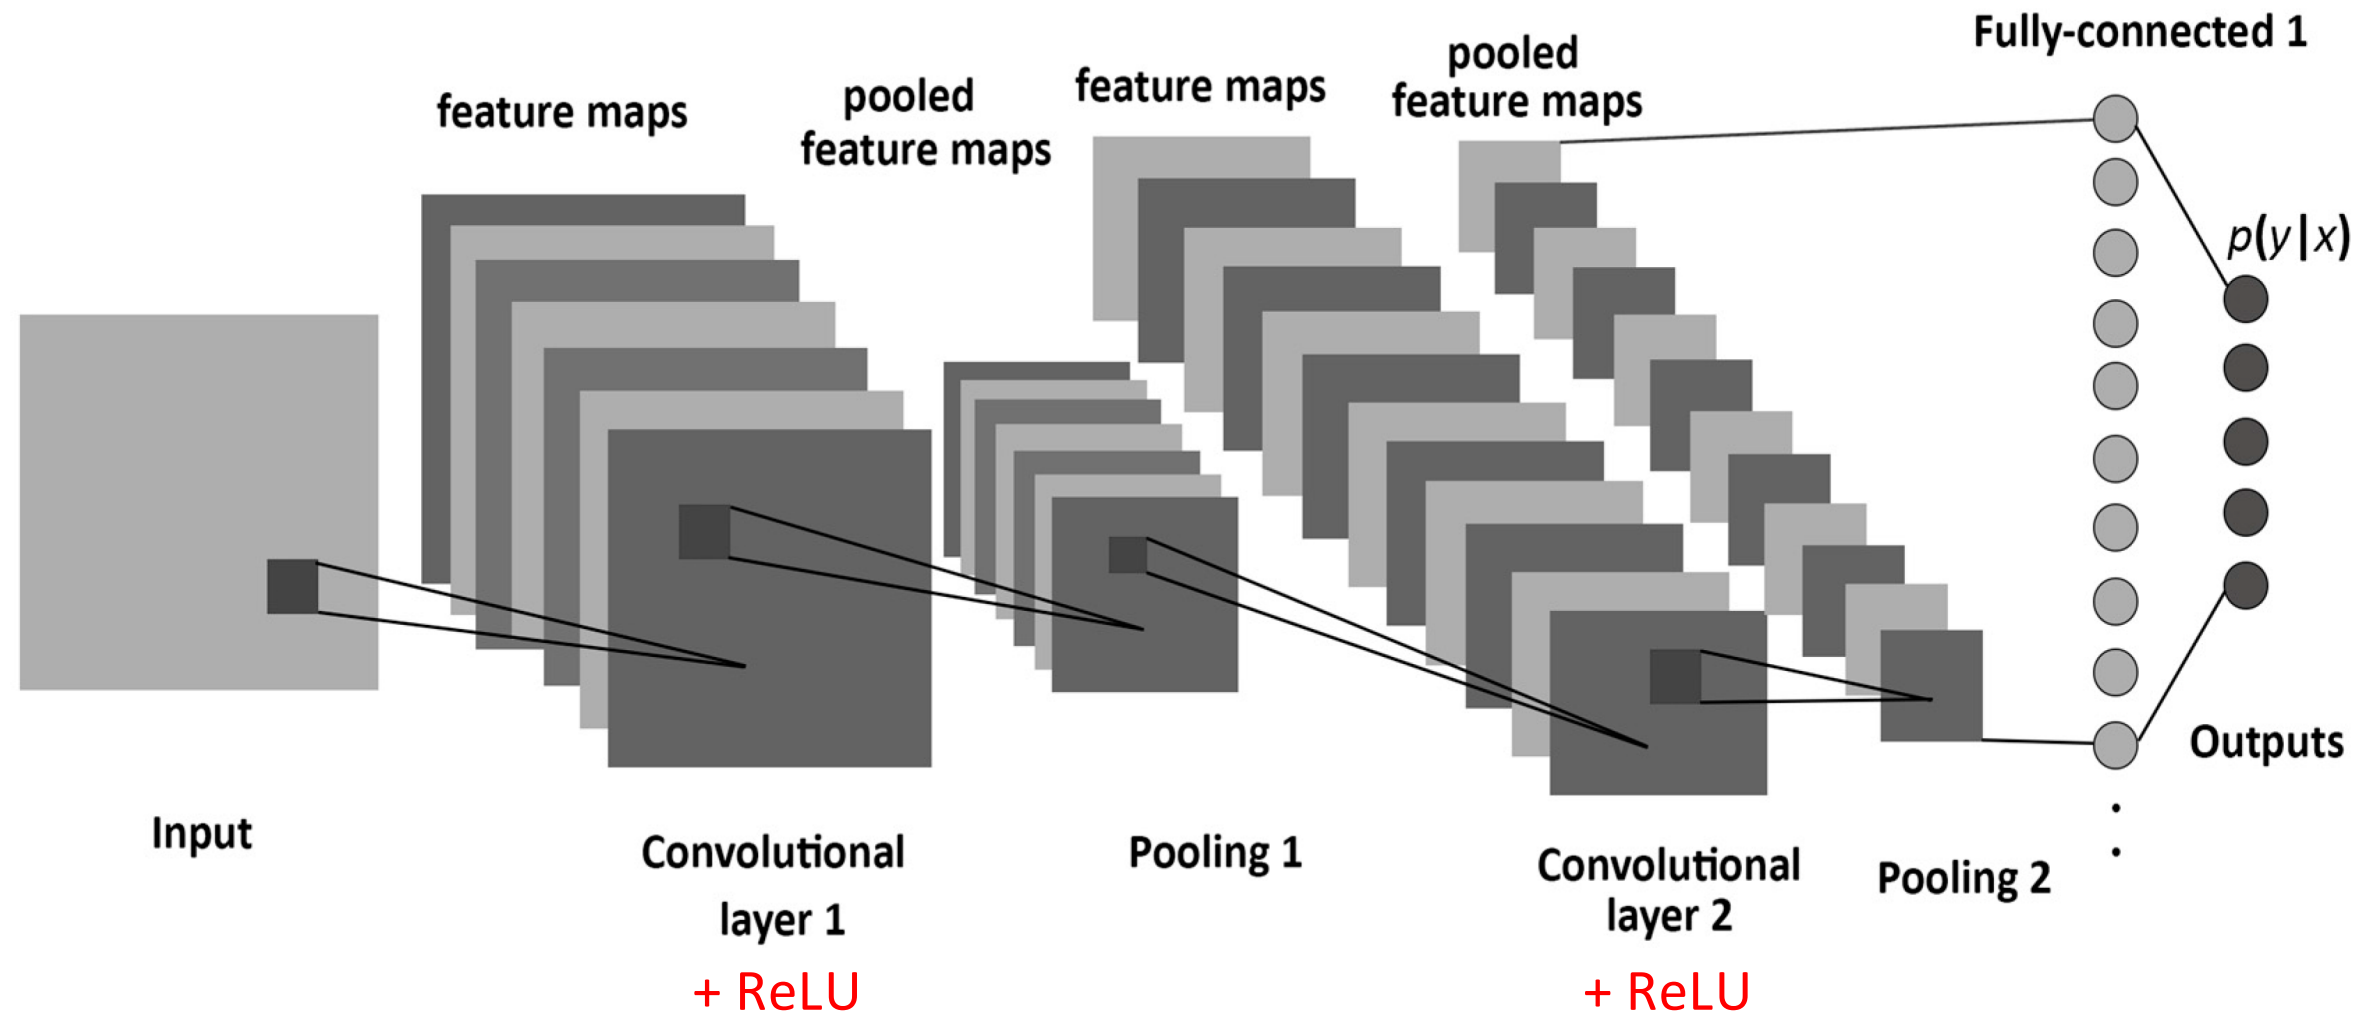
The diagram illustrates the edge detection process. It shows an original image of a squirrel head being convolved with a 3x3 kernel. The kernel is a Laplacian operator, represented by the matrix shown. The result is an edge detection image where only the boundaries of the squirrel are highlighted in white against a black background. An arrow points from the word 'Kernel' to the matrix.

*Sharpen*


$$* \begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix} =$$


The diagram illustrates the sharpening process. It shows the same original image of a squirrel head being convolved with a different 3x3 kernel. This kernel is a sharpening kernel, represented by the matrix shown. The result is a sharpened image where the edges of the squirrel are more pronounced and the overall image appears crisper.


# A convolutional network



# Some key features


- **Sparse interactions:** neurons in a given layer are sensitive only to a sub-field of nearby neurons in the previous layer.
- **Parameter sharing:** same kernel applied to different regions of the input image. Learn kernel only.
- **Pooling:** reduce spatial dimensions of input volume, by averaging over small square windows of the field.
- **Padding:** add additional layer to the border of the image to avoid loss of information.

# Sparse interactions



- Exploits local spatial correlations.
- Receptive field is limited, but increases with depth.
- Reduces memory and run time, as the kernel size is a couple of order of magnitudes smaller than the input image


# Parameter sharing



- Sharing weights helps in translational invariance and also reduces computation time, as there are less free parameters




# Pooling



- Reduced input size without loosing capability to generalize.
- Invariant to small local transitions.

# Padding

- Shift the kernel not by 1 pixel, but by  $k$  pixels (stride)



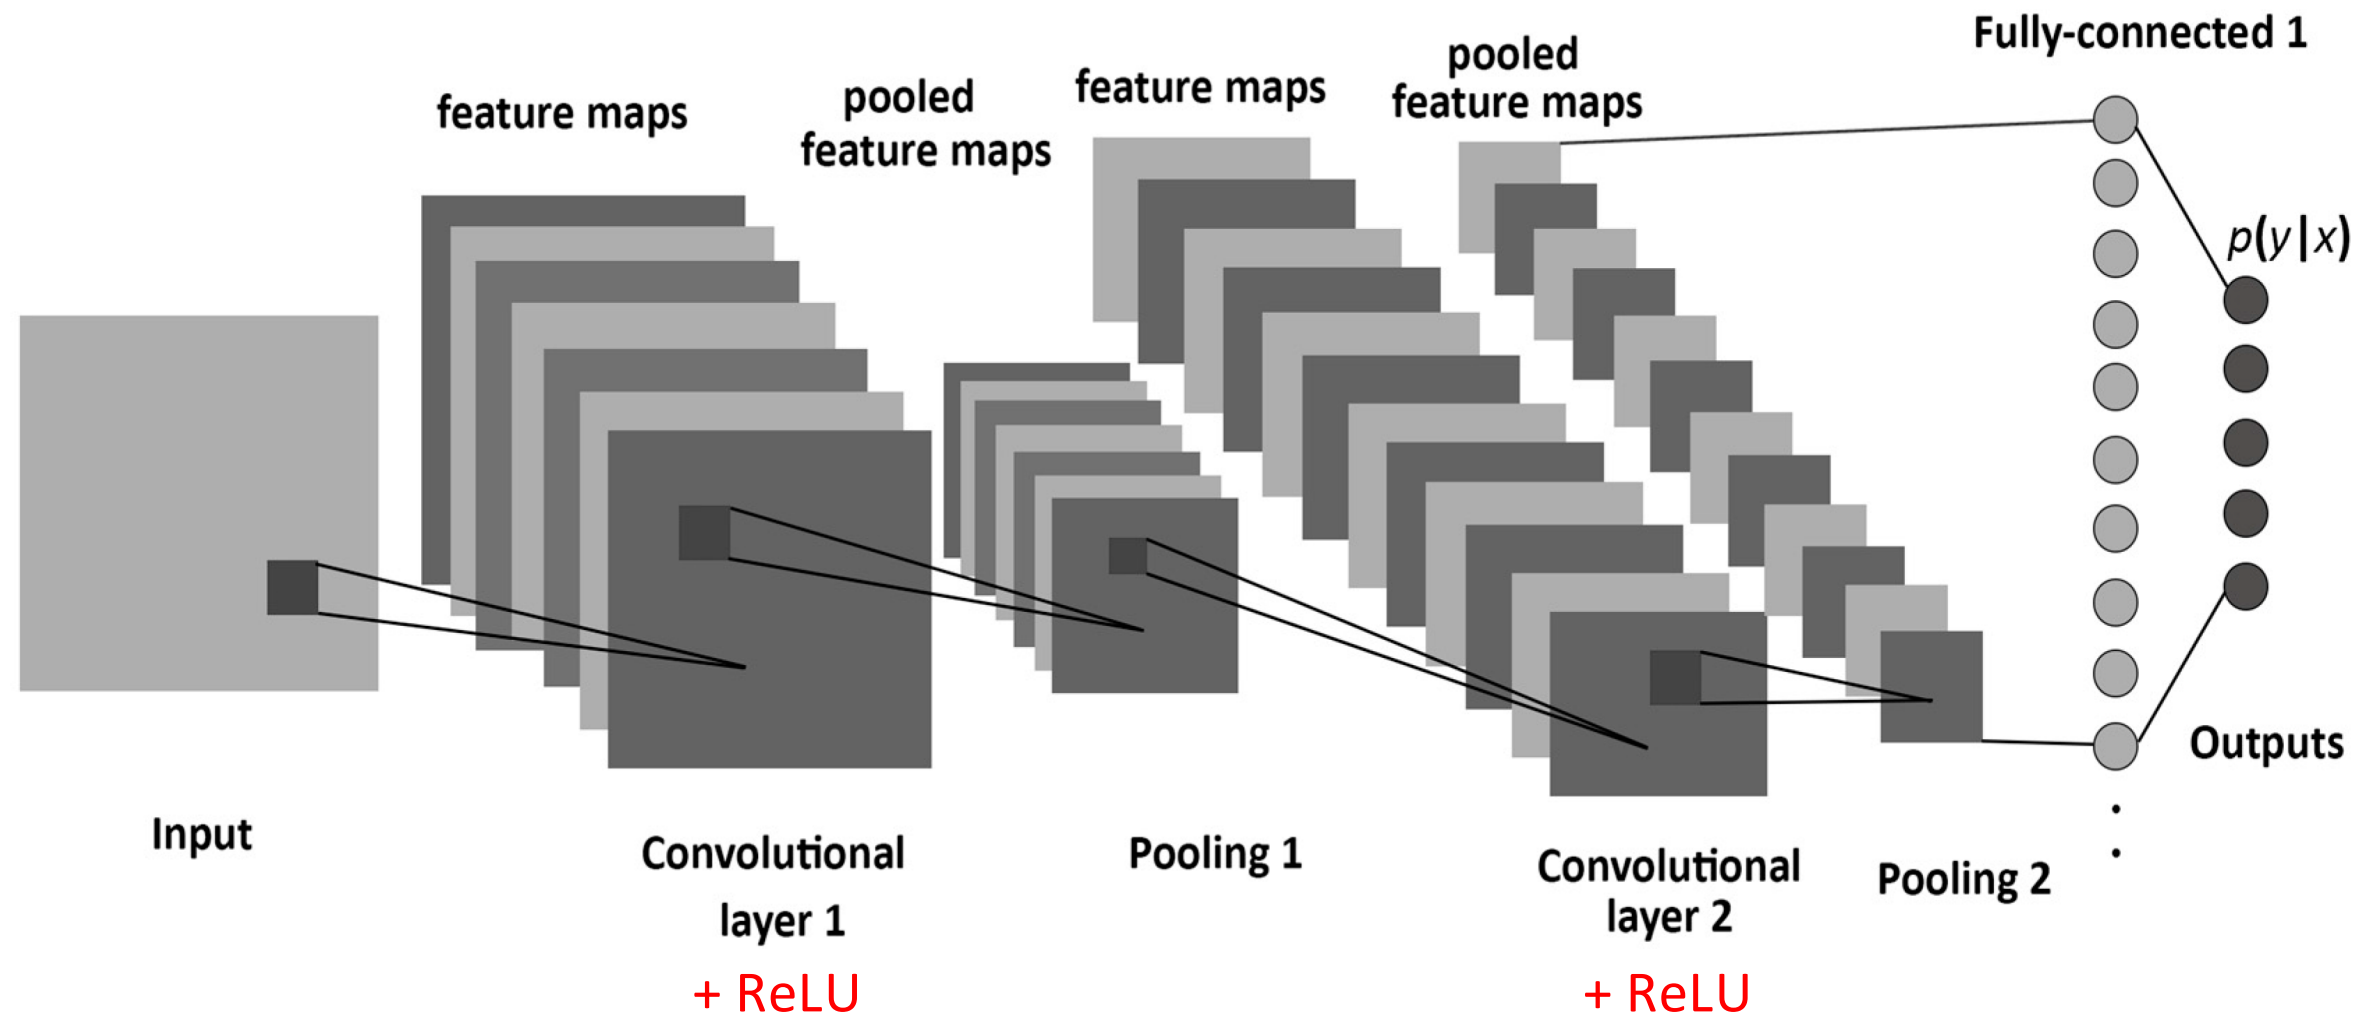
Stride = 2

Sees entire image, but  
low computational and  
storage costs

Prevents loss of information near the edges


ReLU

# A common activation function: ReLU



- Activation functions don't always need to be interpreted as probabilities. They can be anything that produces new representations when applied to the input.
- ReLU: Rectified Linear Unit is a very common activation function used in NNs.
- Defined as:  $y = \max(0, x)$

# A common activation function: ReLU



- Easy to compute: no complicated math.
- Converges fast (gradient always well defined, except for one point).
- Sparsely activated